


Deep dive into coroutines

on the example of Kotlin implementation

Maciej Procyk



March 27, 2023

Table of contents

1 Introduction

- Main ideas
- Practical background

2 Coroutines in details

- Implementation details
- Fundamental concepts
- Structured concurrency

3 Higher-level APIs

- Main ideas
- Practical example

Introduction

General description

- described as “functions whose execution you can pause”

General description

- described as “functions whose execution you can pause”
- computation can be paused and resumed without blocking a thread

General description

- described as “functions whose execution you can pause”
- computation can be paused and resumed without blocking a thread
- generalization of subroutines for cooperative multitasking

General description

- described as “functions whose execution you can pause”
- computation can be paused and resumed without blocking a thread
- generalization of subroutines for cooperative multitasking
- fast context switching

General description

- described as “functions whose execution you can pause”
- computation can be paused and resumed without blocking a thread
- generalization of subroutines for cooperative multitasking
- fast context switching
- perfect tool for implementing iterators, infinite lists, pipes

General description

- described as “functions whose execution you can pause”
- computation can be paused and resumed without blocking a thread
- generalization of subroutines for cooperative multitasking
- fast context switching
- perfect tool for implementing iterators, infinite lists, pipes
- asynchronous code in synchronous manner

Fundamental characteristics

- 1 the values of data local to a coroutine persist between successive calls

Fundamental characteristics

- 1 the values of data local to a coroutine persist between successive calls
- 2 the execution of a coroutine is suspended as control leaves it, only to carry on where it left off when control re-enters the coroutine at some later stage

Asynchronous programming

- nowadays, we write the programs that wait for a lot of time

Asynchronous programming

- nowadays, we write the programs that wait for a lot of time
- we write them using threads, but there are a few limitations

Asynchronous programming

- nowadays, we write the programs that wait for a lot of time
- we write them using threads, but there are a few limitations
- the number of threads cannot be too high, as they consume a lot of memory

Asynchronous programming

- nowadays, we write the programs that wait for a lot of time
- we write them using threads, but there are a few limitations
- the number of threads cannot be too high, as they consume a lot of memory
- e.g. on Linux x64 default stack size (configured with `-XX:ThreadStackSize` or `-Xss` option) is 256 KB

Asynchronous programming

- nowadays, we write the programs that wait for a lot of time
- we write them using threads, but there are a few limitations
- the number of threads cannot be too high, as they consume a lot of memory
- e.g. on Linux x64 default stack size (configured with `-XX:ThreadStackSize` or `-Xss` option) is 256 KB
- which means that 100K threads would consume about 24 GB of memory

Asynchronous programming

- nowadays, we write the programs that wait for a lot of time
- we write them using threads, but there are a few limitations
- the number of threads cannot be too high, as they consume a lot of memory
- e.g. on Linux x64 default stack size (configured with `-XX:ThreadStackSize` or `-Xss` option) is 256 KB
- which means that 100K threads would consume about 24 GB of memory
- that's a lot and what if we think about 1M threads

Coroutines vs threads

- the difference is that threads are scheduled on cores preemptively, while coroutines are scheduled onto threads cooperatively

Coroutines vs threads

- the difference is that threads are scheduled on cores preemptively, while coroutines are scheduled onto threads cooperatively
- starting a thread on JVM requires starting native thread, which consumes memory for its stack

Coroutines vs threads

- the difference is that threads are scheduled on cores preemptively, while coroutines are scheduled onto threads cooperatively
- starting a thread on JVM requires starting native thread, which consumes memory for its stack
- switching between threads requires going through system's kernel

Coroutines vs threads

- the difference is that threads are scheduled on cores preemptively, while coroutines are scheduled onto threads cooperatively
- starting a thread on JVM requires starting native thread, which consumes memory for its stack
- switching between threads requires going through system's kernel
- and is expensive in terms of CPU cycles consumed during that operation

Coroutines vs threads

- coroutine is user-level abstraction (implemented by the language compiler in case of Kotlin)

Coroutines vs threads

- coroutine is user-level abstraction (implemented by the language compiler in case of Kotlin)
- in the simplest case it's a single reference object in JVM heap memory

Coroutines vs threads

- coroutine is user-level abstraction (implemented by the language compiler in case of Kotlin)
- in the simplest case it's a single reference object in JVM heap memory
- switching between coroutines is as simple (and cheap) as invoking a regular function

Sample code starting 100 000 threads

```
1 fun main() {  
2     val counter = AtomicInteger(0)  
3     List(100_000) {  
4         thread {  
5             sleep(1_000)  
6             counter.incrementAndGet()  
7         }  
8     }.forEach { it.join() }  
9     println("counter = ${counter.get()}")  
10 }
```

Sample code starting 100 000 threads

```
1 fun main() {  
2     val counter = AtomicInteger(0)  
3     List(100_000) {  
4         thread {  
5             sleep(1_000)  
6             counter.incrementAndGet()  
7         }  
8     }.forEach { it.join() }  
9     println("counter = ${counter.get()}")  
10 }
```

- possible, but takes a lot of time (and actually they don't exist in the same time)

Sample code starting 100 000 threads

```
1 fun main() {  
2     val counter = AtomicInteger(0)  
3     List(100_000) {  
4         thread {  
5             sleep(1_000)  
6             counter.incrementAndGet()  
7         }  
8     }.forEach { it.join() }  
9     println("counter = ${counter.get()}")  
10 }
```

- possible, but takes a lot of time (and actually they don't exist in the same time)
- if we increase the sleep time the number of threads existing in the same time would increase

Sample code starting 100 000 threads

```
1 fun main() {  
2     val counter = AtomicInteger(0)  
3     List(100_000) {  
4         thread {  
5             sleep(1_000)  
6             counter.incrementAndGet()  
7         }  
8     }.forEach { it.join() }  
9     println("counter = ${counter.get()}")  
10 }
```

- possible, but takes a lot of time (and actually they don't exist in the same time)
- if we increase the sleep time the number of threads existing in the same time would increase
- and we would end up with Exception in thread "main"
java.lang.OutOfMemoryError ...

Sample code starting 100 000 jobs

```
1 fun main() = runBlocking {  
2     val counter = AtomicInteger(0)  
3     List(100_000) {  
4         launch {  
5             delay(1_000)  
6             counter.incrementAndGet()  
7         }  
8     }.forEach { it.join() }  
9     println("counter = ${counter.get()}")  
10 }
```

Sample code starting 100 000 jobs

```
1 fun main() = runBlocking {  
2     val counter = AtomicInteger(0)  
3     List(100_000) {  
4         launch {  
5             delay(1_000)  
6             counter.incrementAndGet()  
7         }  
8     }.forEach { it.join() }  
9     println("counter = ${counter.get()}")  
10 }
```

- explicitly defined context is needed to start the whole process

Sample code starting 100 000 jobs

```
1 fun main() = runBlocking {  
2     val counter = AtomicInteger(0)  
3     List(100_000) {  
4         launch {  
5             delay(1_000)  
6             counter.incrementAndGet()  
7         }  
8     }.forEach { it.join() }  
9     println("counter = ${counter.get()}")  
10 }
```

- explicitly defined context is needed to start the whole process
- code runs over 10 times faster than the threads implementation

Sample code starting 100 000 jobs

```
1 fun main() = runBlocking {  
2     val counter = AtomicInteger(0)  
3     List(100_000) {  
4         launch {  
5             delay(1_000)  
6             counter.incrementAndGet()  
7         }  
8     }.forEach { it.join() }  
9     println("counter = ${counter.get()}")  
10 }
```

- explicitly defined context is needed to start the whole process
- code runs over 10 times faster than the threads implementation
- because it uses computer resources more thoughtful

Toy problem wasting computer resources on sleeping

```
1 fun processRequest() {  
2     val request = receiveRequest()  
3     val processed = handleRequest(request)  
4     sendResponse(processed)  
5 }
```

Toy problem wasting computer resources on sleeping

```
1 fun processRequest() {  
2     val request = receiveRequest()  
3     val processed = handleRequest(request)  
4     sendResponse(processed)  
5 }
```

- potentially, the application may have thousands of these tasks run concurrently

Toy problem wasting computer resources on sleeping

```
1 fun processRequest() {  
2     val request = receiveRequest()  
3     val processed = handleRequest(request)  
4     sendResponse(processed)  
5 }
```

- potentially, the application may have thousands of these tasks run concurrently
- but every of them would consume a single thread to be executed

Toy problem wasting computer resources on sleeping

```
1 fun processRequest() {  
2     val request = receiveRequest()  
3     val processed = handleRequest(request)  
4     sendResponse(processed)  
5 }
```

- potentially, the application may have thousands of these tasks run concurrently
- but every of them would consume a single thread to be executed
- the problem are the blocking functions which are related to IO operations and expensive processing

The “old” JVM solution for this problem are callbacks

```
1 fun processRequestCallbacks() {  
2     receiveRequest { request ->  
3         handleRequest(request) { processed ->  
4             sendResponse(processed)  
5         }  
6     }  
7 }
```

The “old” JVM solution for this problem are callbacks

```
1 fun processRequestCallbacks() {  
2     receiveRequest { request ->  
3         handleRequest(request) { processed ->  
4             sendResponse(processed)  
5         }  
6     }  
7 }
```

- functions returns immediately and run the callbacks later to get the result

The “old” JVM solution for this problem are callbacks

```
1 fun processRequestCallbacks() {  
2     receiveRequest { request ->  
3         handleRequest(request) { processed ->  
4             sendResponse(processed)  
5         }  
6     }  
7 }
```

- functions returns immediately and run the callbacks later to get the result
- this introduces “callback hell” problem

The “old” JVM solution for this problem are callbacks

```
1 fun processRequestCallbacks() {  
2     receiveRequest { request ->  
3         handleRequest(request) { processed ->  
4             sendResponse(processed)  
5         }  
6     }  
7 }
```

- functions returns immediately and run the callbacks later to get the result
- this introduces “callback hell” problem
- this code has multiple limitations - the structure of code prevents imperative coding (how to write loops?)

Writing imperative code with callbacks

```
1 fun imperative() {  
2     var number = 0  
3     for (i in 1..1_000)  
4         number = inc(number)  
5     println(number)  
6 }
```

Writing imperative code with callbacks

```
1 fun imperativeCallback() {  
2     class LateInitCallbackWrapper<T : Any> {  
3         lateinit var callback: (T) -> Unit  
4     }  
5  
6     val calls = Array(1_000) { LateInitCallbackWrapper<Int>() }  
7     for (i in 0..998)  
8         calls[i].callback = { incWithCallback(it, calls[i + 1].callback::invoke) }  
9     calls[999].callback = { incWithCallback(it, ::println) }  
10    calls[0].callback(0)  
11 }
```

Writing imperative code with callbacks

```
1 fun imperativeCallback() {  
2     class LateInitCallbackWrapper<T : Any> {  
3         lateinit var callback: (T) -> Unit  
4     }  
5  
6     val calls = Array(1_000) { LateInitCallbackWrapper<Int>() }  
7     for (i in 0..998)  
8         calls[i].callback = { incWithCallback(it, calls[i + 1].callback::invoke) }  
9     calls[999].callback = { incWithCallback(it, ::println) }  
10    calls[0].callback(0)  
11 }
```

- hard to understand intention of programmer using callbacks

Writing imperative code with callbacks

```
1 fun imperativeCallback() {  
2     class LateInitCallbackWrapper<T : Any> {  
3         lateinit var callback: (T) -> Unit  
4     }  
5  
6     val calls = Array(1_000) { LateInitCallbackWrapper<Int>() }  
7     for (i in 0..998)  
8         calls[i].callback = { incWithCallback(it, calls[i + 1].callback::invoke) }  
9     calls[999].callback = { incWithCallback(it, ::println) }  
10    calls[0].callback(0)  
11 }
```

- hard to understand intention of programmer using callbacks
- additionally, there's extra cost related to working with the references for callbacks

Writing imperative code with callbacks

```
1 fun imperativeCallback() {  
2     class LateInitCallbackWrapper<T : Any> {  
3         lateinit var callback: (T) -> Unit  
4     }  
5  
6     val calls = Array(1_000) { LateInitCallbackWrapper<Int>() }  
7     for (i in 0..998)  
8         calls[i].callback = { incWithCallback(it, calls[i + 1].callback::invoke) }  
9     calls[999].callback = { incWithCallback(it, ::println) }  
10    calls[0].callback(0)  
11 }
```

- hard to understand intention of programmer using callbacks
- additionally, there's extra cost related to working with the references for callbacks
- this solution doesn't scale well so other approach needs to be found

Futures/Promises/Rx being a rescue?

```
1 fun processRequestPromise() {  
2     receiveRequestPromise()  
3     .flatMap { handleRequestPromise(it) }  
4     .subscribeBy(  
5         onSuccess = { sendResponse(it) },  
6         onError = { println("error: $it") },  
7     )  
8 }
```

Futures/Promises/Rx being a rescue?

```
1 fun processRequestPromise() {  
2     receiveRequestPromise()  
3         .flatMap { handleRequestPromise(it) }  
4         .subscribeBy(  
5             onSuccess = { sendResponse(it) },  
6             onError = { println("error: $it") },  
7         )  
8 }
```

- we got nicer, composable code, with no extra indentation

Futures/Promises/Rx being a rescue?

```
1 fun processRequestPromise() {  
2     receiveRequestPromise()  
3         .flatMap { handleRequestPromise(it) }  
4         .subscribeBy(  
5             onSuccess = { sendResponse(it) },  
6             onError = { println("error: $it") },  
7         )  
8 }
```

- we got nicer, composable code, with no extra indentation
- we're able to propagate our exceptions

Futures/Promises/Rx being a rescue?

```
1 fun processRequestPromise() {  
2     receiveRequestPromise()  
3     .flatMap { handleRequestPromise(it) }  
4     .subscribeBy(  
5         onSuccess = { sendResponse(it) },  
6         onError = { println("error: $it") },  
7     )  
8 }
```

- we got nicer, composable code, with no extra indentation
- we're able to propagate our exceptions
- but the combinators are the problem - it might be hard to learn them all

Futures/Promises/Rx being a rescue?

```
1 fun processRequestPromise() {  
2     receiveRequestPromise()  
3     .flatMap { handleRequestPromise(it) }  
4     .subscribeBy(  
5         onSuccess = { sendResponse(it) },  
6         onError = { println("error: $it") },  
7     )  
8 }
```

- we got nicer, composable code, with no extra indentation
- we're able to propagate our exceptions
- but the combinators are the problem - it might be hard to learn them all
- while there are different libraries with other names

Kotlin solution are coroutines

```
1 suspend fun processRequestCoroutines() {  
2     val request = receiveRequestCoroutines()  
3     val processed = handleRequestCoroutines(request)  
4     sendResponseCoroutines(processed)  
5 }
```

Kotlin solution are coroutines

```
1 suspend fun processRequestCoroutines() {  
2     val request = receiveRequestCoroutines()  
3     val processed = handleRequestCoroutines(request)  
4     sendResponseCoroutines(processed)  
5 }
```

- need to mark functions with `suspend` modifier

Kotlin solution are coroutines

```
1 suspend fun processRequestCoroutines() {  
2     val request = receiveRequestCoroutines()  
3     val processed = handleRequestCoroutines(request)  
4     sendResponseCoroutines(processed)  
5 }
```

- need to mark functions with `suspend` modifier
- but the structure of the code looks like regular code (loops, exceptions, extension functions)

Kotlin solution are coroutines

```
1 suspend fun processRequestCoroutines() {  
2     val request = receiveRequestCoroutines()  
3     val processed = handleRequestCoroutines(request)  
4     sendResponseCoroutines(processed)  
5 }
```

- need to mark functions with `suspend` modifier
- but the structure of the code looks like regular code (loops, exceptions, extension functions)
- we got “not obvious” suspension points

Kotlin solution are coroutines

```
1 suspend fun processRequestCoroutines() {  
2     val request = receiveRequestCoroutines()  
3     val processed = handleRequestCoroutines(request)  
4     sendResponseCoroutines(processed)  
5 }
```

- need to mark functions with `suspend` modifier
- but the structure of the code looks like regular code (loops, exceptions, extension functions)
- we got “not obvious” suspension points
- but the good IDE helps here a lot if we're interested in this

Coroutines in details

What's under the hood?

- suspending functions are implemented via Continuation-Passing-Style (CPS)

What's under the hood?

- suspending functions are implemented via Continuation-Passing-Style (CPS)
- every suspending function has an additional `Continuation` parameter passed implicitly when function is invoked

What's under the hood?

- suspending functions are implemented via Continuation-Passing-Style (CPS)
- every suspending function has an additional `Continuation` parameter passed implicitly when function is invoked
- consider the following signature of `suspend` function

What's under the hood?

- suspending functions are implemented via Continuation-Passing-Style (CPS)
- every suspending function has an additional `Continuation` parameter passed implicitly when function is invoked
- consider the following signature of `suspend` function

```
suspend fun <T> CompletableFuture<T>.await(): T =
```

What's under the hood?

- suspending functions are implemented via Continuation-Passing-Style (CPS)
- every suspending function has an additional `Continuation` parameter passed implicitly when function is invoked
- consider the following signature of `suspend` function

```
suspend fun <T> CompletableFuture<T>.await(): T =
```

- its actual implementation after compiler transformation has the following form

What's under the hood?

- suspending functions are implemented via Continuation-Passing-Style (CPS)
- every suspending function has an additional `Continuation` parameter passed implicitly when function is invoked
- consider the following signature of `suspend` function

```
suspend fun <T> CompletableFuture<T>.await(): T =
```

- its actual implementation after compiler transformation has the following form

```
fun <T> CompletableFuture<T>.await(continuation: Continuation<T>): Any? =
```

Continuation interface

- context represents an arbitrary user-defined context that is associated with the coroutine

```
1 interface Continuation<in T> {  
2     val context: CoroutineContext  
3     fun resumeWith(result: Result<T>)  
4 }
```

Continuation interface

- `context` represents an arbitrary user-defined context that is associated with the coroutine
- `resumeWith` function is a completion callback that is used to report on coroutine completion:

```
1 interface Continuation<in T> {  
2     val context: CoroutineContext  
3     fun resumeWith(result: Result<T>)  
4 }
```


Continuation interface

- `context` represents an arbitrary user-defined context that is associated with the coroutine
- `resumeWith` function is a completion callback that is used to report on coroutine completion:
 - a success - with a value

```
1 interface Continuation<in T> {  
2     val context: CoroutineContext  
3     fun resumeWith(result: Result<T>)  
4 }
```

Continuation interface

- `context` represents an arbitrary user-defined context that is associated with the coroutine
- `resumeWith` function is a completion callback that is used to report on coroutine completion:
 - a success - with a value
 - a failure - with an exception

```
1 interface Continuation<in T> {  
2     val context: CoroutineContext  
3     fun resumeWith(result: Result<T>)  
4 }
```

Continuation interface

- `context` represents an arbitrary user-defined context that is associated with the coroutine
- `resumeWith` function is a completion callback that is used to report on coroutine completion:
 - a success - with a value
 - a failure - with an exception
- there are two extension functions defined for convenience

```
1 interface Continuation<in T> {  
2     val context: CoroutineContext  
3     fun resumeWith(result: Result<T>)  
4 }
```

```
1 fun <T> Continuation<T>.resume(value: T) { .. }  
2  
3 fun <T> Continuation<T>.resumeWithException(exception: Throwable) { .. }
```

How the typical implementation looks like?

```
1 suspend fun <T> CompletableFuture<T>.await(): T =  
2     suspendCoroutine { cont: Continuation<T> ->  
3         whenComplete { result, exception ->  
4             if (exception == null) cont.resume(result)  
5             else cont.resumeWithException(exception)  
6         }  
7     }
```

How the typical implementation looks like?

```
1 suspend fun <T> CompletableFuture<T>.await(): T =
2     suspendCoroutine { cont: Continuation<T> ->
3         (this as CompletableFuture<T>).whenComplete(
4             object : BiConsumer<T, Throwable?> {
5                 override fun accept(result: T, exception: Throwable?) {
6                     if (exception == null) cont.resume(result)
7                     else cont.resumeWithException(exception)
8                 }
9             }
10        )
11    }
```

How the typical implementation looks like?

```
1 suspend fun <T> CompletableFuture<T>.await(): T =
2     suspendCoroutine { cont: Continuation<T> ->
3         (this as CompletableFuture<T>).whenComplete(
4             object : BiConsumer<T, Throwable?> {
5                 override fun accept(result: T, exception: Throwable?) {
6                     if (exception == null) cont.resume(result)
7                     else cont.resumeWithException(exception)
8                 }
9             }
10        )
11    }
```

- we use the callback interface to resume the continuation

How the typical implementation looks like?

```

1  suspend fun <T> CompletableFuture<T>.await(): T =
2      suspendCoroutine { cont: Continuation<T> ->
3          (this as CompletableFuture<T>).whenComplete(
4              object : BiConsumer<T, Throwable?> {
5                  override fun accept(result: T, exception: Throwable?) {
6                      if (exception == null) cont.resume(result)
7                      else cont.resumeWithException(exception)
8                  }
9              }
10         )
11     }

```

- we use the callback interface to resume the continuation
- to actually suspend execution function must invoke other suspending function - await implementation invokes a suspending function `suspendCoroutine`

State machines as perfect solution

- it's crucial to implement coroutines efficiently, i.e. create as few classes and objects as possible

State machines as perfect solution

- it's crucial to implement coroutines efficiently, i.e. create as few classes and objects as possible
- compiler creates a single instance of class that may have any number of suspension points

State machines as perfect solution

- it's crucial to implement coroutines efficiently, i.e. create as few classes and objects as possible
- compiler creates a single instance of class that may have any number of suspension points
- general idea - suspending function is a state machine

State machines as perfect solution

- it's crucial to implement coroutines efficiently, i.e. create as few classes and objects as possible
- compiler creates a single instance of class that may have any number of suspension points
- general idea - suspending function is a state machine
- states corresponds to suspension points

State machines as perfect solution

- it's crucial to implement coroutines efficiently, i.e. create as few classes and objects as possible
- compiler creates a single instance of class that may have any number of suspension points
- general idea - suspending function is a state machine
- states corresponds to suspension points

```
1 suspend fun stateMachine() {  
2     val a = a()  
3     val y = foo(a).await() // 1st suspension point  
4     b()  
5     val z = bar(a, y).await() // 2nd suspension point  
6     c(z)  
7 }
```

State machines as perfect solution

- three states for the given block of code:

State machines as perfect solution

- three states for the given block of code:
 - 1 initial (before any suspension point)

State machines as perfect solution

- three states for the given block of code:
 - 1 initial (before any suspension point)
 - 2 after the first suspension point

State machines as perfect solution

- three states for the given block of code:
 - 1 initial (before any suspension point)
 - 2 after the first suspension point
 - 3 after the second suspension point

State machines as perfect solution

- three states for the given block of code:
 - 1 initial (before any suspension point)
 - 2 after the first suspension point
 - 3 after the second suspension point
- each of them being an entry point to continuation of this block

State machines as perfect solution

- three states for the given block of code:
 - 1 initial (before any suspension point)
 - 2 after the first suspension point
 - 3 after the second suspension point
- each of them being an entry point to continuation of this block
- which can be seen as `switch` statement with label to each suspension point

State machines as perfect solution

- three states for the given block of code:
 - 1 initial (before any suspension point)
 - 2 after the first suspension point
 - 3 after the second suspension point
- each of them being an entry point to continuation of this block
- which can be seen as `switch` statement with label to each suspension point
- the execution is continued with `resumeWith` method of coroutine which selects the current branch

State machines as perfect solution

- three states for the given block of code:
 - 1 initial (before any suspension point)
 - 2 after the first suspension point
 - 3 after the second suspension point
- each of them being an entry point to continuation of this block
- which can be seen as `switch` statement with label to each suspension point
- the execution is continued with `resumeWith` method of coroutine which selects the current branch
- local variables are generated as class fields of anonymous class

State machines implementation

Pseudo-code corresponding to JVM bytecode created for suspending function:

```
1  class <anonymous_for_state_machine> extends SuspendLambda<...> {
2      // current state of the state machine
3      int label = 0
4
5      // local variables of the coroutine
6      A a = null
7      Y y = null
8
9      void resumeWith(Object result) {
10         if (label == 0) goto L0
11         if (label == 1) goto L1
12         if (label == 2) goto L2
13         else throw IllegalStateException()
14         // each label code
```

State machines implementation

```
15         // local variables and state switch
16     L0:
17         // result is expected to be `null` at this invocation
18         a = a()
19         label = 1
20         result = foo(a).await(this) // 'this' is passed as a continuation
21         if (result == COROUTINE_SUSPENDED) return // return if await had suspended
22     L1:
23         // external code has resumed this coroutine passing the result of .await()
24         y = (Y) result
25         b()
26         label = 2
27         result = bar(a, y).await(this) // 'this' is passed as a continuation
28         if (result == COROUTINE_SUSPENDED) return // return if await had suspended
29     L2:
30         // external code has resumed this coroutine passing the result of .await()
31         z = (Z) result
32         c(z)
33         label = -1 // No more steps are allowed
34         return
35     }
36 }
```

Generators

- generators with `yield` can be easily implemented using suspend functions

Generators

- generators with `yield` can be easily implemented using suspend functions
- and that's how the Kotlin std-lib does, providing `Sequence<T>` abstraction with `yield` as a function of `SequenceScope<T>`

Generators

- generators with `yield` can be easily implemented using suspend functions
- and that's how the Kotlin std-lib does, providing `Sequence<T>` abstraction with `yield` as a function of `SequenceScope<T>`
- they can represent potentially infinite lists (like in Haskell)

Generators

- generators with `yield` can be easily implemented using suspend functions
- and that's how the Kotlin std-lib does, providing `Sequence<T>` abstraction with `yield` as a function of `SequenceScope<T>`
- they can represent potentially infinite lists (like in Haskell)
- or just in general lazily computed sequences

Generators

- generators with `yield` can be easily implemented using suspend functions
- and that's how the Kotlin std-lib does, providing `Sequence<T>` abstraction with `yield` as a function of `SequenceScope<T>`
- they can represent potentially infinite lists (like in Haskell)
- or just in general lazily computed sequences
- the strength of them is supporting arbitrary control flow, with `for`, `when`, `try/catch` etc.

Generators - example Fibonacci sequence

```
1 fun fibonacci() = sequence {  
2     yield(1)  
3     var curr = 1  
4     var next = 1  
5     while (true) {  
6         yield(next)  
7         val tmp = curr + next  
8         curr = next  
9         next = tmp  
10    }  
11 }
```

Generators - example Fibonacci sequence

```
1 fun fibonacci() = sequence {  
2     yield(1)  
3     var curr = 1  
4     var next = 1  
5     while (true) {  
6         yield(next)  
7         val tmp = curr + next  
8         curr = next  
9         next = tmp  
10    }  
11 }
```

```
fun main() = fibonacci().take(42).forEach(::println)
```

Initial idea - extra DSL

```
1 fun processRequest() = async {  
2     val request = await(receiveRequestPromise())  
3     val processed = await(handleRequestPromise(request))  
4     sendResponse(processed)  
5 }
```

Initial idea - extra DSL

```
1 fun processRequest() = async {  
2     val request = await(receiveRequestPromise())  
3     val processed = await(handleRequestPromise(request))  
4     sendResponse(processed)  
5 }
```

- original prototype defined the library functions, similar to keywords from e.g. C#, Typescript, Dart, Python

Initial idea - extra DSL

```
1 fun processRequest() = async {  
2     val request = await(receiveRequestPromise())  
3     val processed = await(handleRequestPromise(request))  
4     sendResponse(processed)  
5 }
```

- original prototype defined the library functions, similar to keywords from e.g. C#, Typescript, Dart, Python
- but the compiler needed to know that `await` can be suspended, so it was marked with `suspend` modifier

Initial idea - extra DSL

```
1 fun processRequest() = async {  
2     val request = await(receiveRequestPromise())  
3     val processed = await(handleRequestPromise(request))  
4     sendResponse(processed)  
5 }
```

- original prototype defined the library functions, similar to keywords from e.g. C#, Typescript, Dart, Python
- but the compiler needed to know that `await` can be suspended, so it was marked with `suspend` modifier
- so the decision was to remove the need for `await` call

Initial idea - extra DSL

```
1 fun processRequest() = async {  
2     val request = await(receiveRequestPromise())  
3     val processed = await(handleRequestPromise(request))  
4     sendResponse(processed)  
5 }
```

- original prototype defined the library functions, similar to keywords from e.g. C#, Typescript, Dart, Python
- but the compiler needed to know that `await` can be suspended, so it was marked with `suspend` modifier
- so the decision was to remove the need for `await` call
- which led to adding `suspend` modifier to functions instead of using `async` builder

Initial idea - extra DSL

- there was some fear as no other language resigned from `async/await`

Initial idea - extra DSL

- there was some fear as no other language resigned from `async/await`
- actually in Go language nobody is worried about the not visible suspension points

Initial idea - extra DSL

- there was some fear as no other language resigned from `async/await`
- actually in Go language nobody is worried about the not visible suspension points

```
1 suspend fun go() = coroutineScope {  
2     go { printDelayed("go: fizz") }  
3     printDelayed("go: buzz")  
4 }
```

Initial idea - extra DSL

- there was some fear as no other language resigned from `async/await`
- actually in Go language nobody is worried about the not visible suspension points

```
1 suspend fun go() = coroutineScope {  
2     go { printDelayed("go: fizz") }  
3     printDelayed("go: buzz")  
4 }
```

```
1 suspend fun kt() = coroutineScope {  
2     launch { printDelayed("kt: fizz") }  
3     printDelayed("kt: buzz")  
4 }
```

Initial idea - extra DSL

- there was some fear as no other language resigned from `async/await`
- actually in Go language nobody is worried about the not visible suspension points

```
1 suspend fun go() = coroutineScope {
2     go { printDelayed("go: fizz") }
3     printDelayed("go: buzz")
4 }
```

```
1 suspend fun kt() = coroutineScope {
2     launch { printDelayed("kt: fizz") }
3     printDelayed("kt: buzz")
4 }
```

```
1 suspend fun printDelayed(msg: Any, times: Int = 5) {
2     for (i in 1..times) {
3         delay(1_000)
4         println(msg)
5     }
6 }
```

Context and Scope of execution

```
1 suspend fun implicitScope(): Unit = coroutineScope {  
2     launch(context = Default) { say("hello") }  
3     launch(context = IO) { say("world") }  
4 }
```


Context and Scope of execution

```
1 suspend fun implicitScope(): Unit = coroutineScope {  
2     launch(context = Default) { say("hello") }  
3     launch(context = IO) { say("world") }  
4 }
```

- coroutines execute in some context represented by the `CoroutineContext` type

Context and Scope of execution

```
1 suspend fun implicitScope(): Unit = coroutineScope {  
2     launch(context = Default) { say("hello") }  
3     launch(context = IO) { say("world") }  
4 }
```

- coroutines execute in some context represented by the `CoroutineContext` type
- it contains the `CoroutineDispatcher` which delegates the coroutine to some thread of execution

Context and Scope of execution

```
1 suspend fun implicitScope(): Unit = coroutineScope {  
2     launch(context = Default) { say("hello") }  
3     launch(context = IO) { say("world") }  
4 }
```

- coroutines execute in some context represented by the `CoroutineContext` type
- it contains the `CoroutineDispatcher` which delegates the coroutine to some thread of execution
- it's inherited from the parent `CoroutineScope` by default

Context and Scope of execution

```
1 suspend fun implicitScope(): Unit = coroutineScope {  
2     launch(context = Default) { say("hello") }  
3     launch(context = IO) { say("world") }  
4 }
```

- coroutines execute in some context represented by the `CoroutineContext` type
- it contains the `CoroutineDispatcher` which delegates the coroutine to some thread of execution
- it's inherited from the parent `CoroutineScope` by default
- and all coroutine builders like `launch` and `async` accept optional `CoroutineContext` to explicitly specify the dispatcher for the new coroutine context

Context and Scope of execution

```
1 suspend fun explicitScope(): Unit = coroutineScope {  
2     val scope: CoroutineScope = this  
3     scope.launch(context = Default) { say("hello") }  
4     scope.launch(context = IO) { say("world") }  
5 }
```

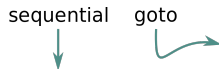
- coroutines execute in some context represented by the `CoroutineContext` type
- it contains the `CoroutineDispatcher` which delegates the coroutine to some thread of execution
- it's inherited from the parent `CoroutineScope` by default
- and all coroutine builders like `launch` and `async` accept optional `CoroutineContext` to explicitly specify the dispatcher for the new coroutine context

What is structured concurrency?

- compare understanding concurrency to understanding program control flow

What is structured concurrency?

- compare understanding concurrency to understanding program control flow
- with goto instruction it can start being hard to understand




source:

<https://vorp.us.org/blog/notes-on-structured-concurrency-or-go-statement-considered-harmful/>

What is structured concurrency?

- compare understanding concurrency to understanding program control flow
- with goto instruction it can start being hard to understand
- but when we limit ourselves to known structures, it's much more obvious and not that limited



source:


<https://vorpus.org/blog/notes-on-structured-concurrency-or-go-statement-considered-harmful/>

What is structured concurrency?

- compare understanding concurrency to understanding program control flow
- with goto instruction it can start being hard to understand
- but when we limit ourselves to known structures, it's much more obvious and not that limited
- in concurrent programming model we have the same problem with tasks

What is structured concurrency?

- compare understanding concurrency to understanding program control flow
- with goto instruction it can start being hard to understand
- but when we limit ourselves to known structures, it's much more obvious and not that limited
- in concurrent programming model we have the same problem with tasks
- introducing some “structures” to keep everything in consistent state helps in understanding what may happen



source:
<https://vorpus.org/blog/notes-on-structured-concurrency-or-go-statement-considered-harmful/>

launch and async coroutine builders

- launch produces a Job, while async produces a Deferred<T>

launch and async coroutine builders

- launch produces a Job, while async produces a Deferred<T>
- they start a new coroutine without blocking the current one

launch and async coroutine builders

- launch produces a Job, while async produces a Deferred<T>
- they start a new coroutine without blocking the current one
- they can be called only in context of some CoroutineScope which is “responsible” for its execution

launch and async coroutine builders

- `launch` produces a `Job`, while `async` produces a `Deferred<T>`
- they start a new coroutine without blocking the current one
- they can be called only in context of some `CoroutineScope` which is “responsible” for its execution
- coroutine is cancelled when the returned `Job` is cancelled

launch and async coroutine builders

- `launch` produces a `Job`, while `async` produces a `Deferred<T>`
- they start a new coroutine without blocking the current one
- they can be called only in context of some `CoroutineScope` which is “responsible” for its execution
- coroutine is cancelled when the returned `Job` is cancelled
- but still we need to manually check for cancellation in long-running tasks

Structured concurrency in Kotlin

- CoroutineScope can be seen a structure in which the control flows of concurrent tasks are merged

Structured concurrency in Kotlin

- CoroutineScope can be seen a structure in which the control flows of concurrent tasks are merged
- launch and async corresponds to creating new execution branches

Structured concurrency in Kotlin

- CoroutineScope can be seen a structure in which the control flows of concurrent tasks are merged
- launch and async corresponds to creating new execution branches
- parent always waits for children completion

Structured concurrency in Kotlin

- CoroutineScope can be seen a structure in which the control flows of concurrent tasks are merged
- launch and async corresponds to creating new execution branches
- parent always waits for children completion
- there's no place for loosing some resource

Structured concurrency in Kotlin

- CoroutineScope can be seen a structure in which the control flows of concurrent tasks are merged
- launch and async corresponds to creating new execution branches
- parent always waits for children completion
- there's no place for losing some resource
- there's no exception lost - they're propagated

launch running examples

```
1 suspend fun justSuspensionPoints(): Unit = coroutineScope {  
2     sayA()  
3     sayB()  
4 }
```

launch running examples

```
1 suspend fun justSuspensionPoints(): Unit = coroutineScope {  
2     sayA()  
3     sayB()  
4 }
```

- executes sayA and sayB sequentially

launch running examples

```
1 suspend fun justSuspensionPoints(): Unit = coroutineScope {  
2     sayA()  
3     sayB()  
4 }
```

- executes sayA and sayB sequentially
- finishes justSuspensionPoints when both jobs are finished

launch running examples

```
1 suspend fun launchFirstInScope(): Unit = coroutineScope {  
2     launch { sayA() }  
3     sayB()  
4 }
```


launch running examples

```
1 suspend fun launchFirstInScope(): Unit = coroutineScope {  
2     launch { sayA() }  
3     sayB()  
4 }
```

- executes sayA and sayB concurrently

launch running examples

```
1 suspend fun launchFirstInScope(): Unit = coroutineScope {  
2     launch { sayA() }  
3     sayB()  
4 }
```

- executes sayA and sayB concurrently
- finishes launchFirstInScope when both jobs are finished

launch running examples

```
1 suspend fun launchBothInScope(): Unit = coroutineScope {  
2     launch { sayA() }  
3     launch { sayB() }  
4 }
```

launch running examples

```
1 suspend fun launchBothInScope(): Unit = coroutineScope {  
2     launch { sayA() }  
3     launch { sayB() }  
4 }
```

- executes sayA and sayB concurrently

launch running examples

```
1 suspend fun launchBothInScope(): Unit = coroutineScope {  
2     launch { sayA() }  
3     launch { sayB() }  
4 }
```

- executes sayA and sayB concurrently
- finishes launchBothInScope when both jobs are finished

launch running examples

```
1 suspend fun launchWithException(): Unit = coroutineScope {  
2     launch { error("illegal to sayA") }  
3     launch { sayB() }  
4 }
```

launch running examples

```
1 suspend fun launchWithException(): Unit = coroutineScope {  
2     launch { error("illegal to sayA") }  
3     launch { sayB() }  
4 }
```

- executes sayA and sayB concurrently

launch running examples

```
1 suspend fun launchWithException(): Unit = coroutineScope {  
2     launch { error("illegal to sayA") }  
3     launch { sayB() }  
4 }
```

- executes sayA and sayB concurrently
- when sayA fails, sayB is cancelled as well

launch running examples

```
1 suspend fun launchInSupervisorWithException(): Unit = supervisorScope {  
2     launch { error("illegal to sayA") }  
3     launch { sayB() }  
4 }
```

launch running examples

```
1 suspend fun launchInSupervisorWithException(): Unit = supervisorScope {  
2     launch { error("illegal to sayA") }  
3     launch { sayB() }  
4 }
```

- executes sayA and sayB concurrently

launch running examples

```
1 suspend fun launchInSupervisorWithException(): Unit = supervisorScope {  
2     launch { error("illegal to sayA") }  
3     launch { sayB() }  
4 }
```

- executes sayA and sayB concurrently
- when sayA fails, sayB is executed until it finishes

Higher-order function

- introducing coroutine brings advantage of no extra need for higher-order combinators like `flatMap`

Higher-order function

- introducing coroutine brings advantage of no extra need for higher-order combinators like `flatMap`
- it's enough to use the `suspend` modifier for lambda parameter

Higher-order function

- introducing coroutine brings advantage of no extra need for higher-order combinators like flatMap
- it's enough to use the suspend modifier for lambda parameter

```
1 suspend fun main() = retry(afterMillis = 10) { say("hello") }
2
3 suspend fun retry(afterMillis: Long, action: suspend () -> Unit) {
4     do {
5         try {
6             return action()
7         } catch (e: Exception) {
8             println(e.stackTrace)
9         }
10        delay(afterMillis)
11    } while (coroutineContext.isActive)
12 }
```

Higher-level APIs

Channel<T> in Kotlin

- is implemented in Kotlin as a library, conceptually very similar to `BlockingQueue`

Channel<T> in Kotlin

- is implemented in Kotlin as a library, conceptually very similar to `BlockingQueue`
- implements `SendChannel<T>` and `ReceiveChannel<T>`

```
1 interface SendChannel<T> {  
2     suspend fun send(value: T)  
3     fun close()  
4 }  
5  
6 interface ReceiveChannel<T> {  
7     suspend fun receive(): T  
8     suspend operator fun iterator(): ReceiveIterator<T>  
9 }
```

Channel<T> in Kotlin

- is implemented in Kotlin as a library, conceptually very similar to `BlockingQueue`
- implements `SendChannel<T>` and `ReceiveChannel<T>`
- `send` suspends when the channel buffer is full, while `receive` suspends when the buffer is empty

Channel<T> in Kotlin

- is implemented in Kotlin as a library, conceptually very similar to `BlockingQueue`
- implements `SendChannel<T>` and `ReceiveChannel<T>`
- `send` suspends when the channel buffer is full, while `receive` suspends when the buffer is empty
- channels are hot - there's a coroutine on the other side of the channel that produces values, so we cannot just drop a reference to the `ReceiveChannel`, because the producer is going to be suspended forever waiting for a consumer, wasting memory resources, open network connections, etc.

Flow<T> in Kotlin

- represents the stream of values that are being computed asynchronously

Flow<T> in Kotlin

- represents the stream of values that are being computed asynchronously
- can be used just like a `Sequence<T>` type for synchronously computed values

Flow<T> in Kotlin

- represents the stream of values that are being computed asynchronously
- can be used just like a `Sequence<T>` type for synchronously computed values
- flows are cold streams similar to sequences

Flow<T> in Kotlin

- represents the stream of values that are being computed asynchronously
- can be used just like a `Sequence<T>` type for synchronously computed values
- flows are cold streams similar to sequences
- the code inside a flow builder does not run until the flow is collected

Images downloader - goals

- have single abstraction responsible for caching the results

Images downloader - goals

- have single abstraction responsible for caching the results
- have multiple workers responsible for doing the job

Images downloader - goals

- have single abstraction responsible for caching the results
- have multiple workers responsible for doing the job
- communicate between these abstractions in safely way

Images downloader - goals

- have single abstraction responsible for caching the results
- have multiple workers responsible for doing the job
- communicate between these abstractions in safely way
- structure our concurrency model to be sure what's going on

Images downloader - downloader

```
1 fun CoroutineScope.downloader(  
2     receiveStringUrlChannel: ReceiveChannel<String>,  
3     sendResultByteArrayChannel: SendChannel<ByteArray>,  
4     sendToWorkerChannel: SendChannel<Url>,  
5     receiveFromWorkerChannel: ReceiveChannel<DownloadedData>,  
6 ) = launch {  
7     val cached = mutableMapOf<Url, ByteArray>()  
8     while (isActive) {  
9         select {  
10             receiveFromWorkerChannel.onReceive { (url, data) ->  
11                 cached[url] = data  
12                 sendResultByteArrayChannel.send(data)  
13             }  
14             receiveStringUrlChannel.onReceive {  
15                 val url = Url(it)  
16                 val data = cached[url]  
17                 if (data == null) sendToWorkerChannel.send(url)  
18                 else sendResultByteArrayChannel.send(data)  
19             }  
20         }  
21     }  
22 }
```

Images downloader - worker

```
1 fun CoroutineScope.downloadWorker(  
2     receive: ReceiveChannel<Url>,  
3     sendData: SendChannel<DownloadedData>,  
4 ) = launch {  
5     val client = HttpClient(CIO)  
6     for (url in receive) {  
7         val data = client.get(url)  
8         val bytes = data.readBytes()  
9         val downloaded = DownloadedData(url, bytes)  
10        sendData.send(downloaded)  
11    }  
12 }
```

Images downloader - common scope


```
1 fun CoroutineScope.processUrls(  
2     receiveStringUrlChannel: ReceiveChannel<String>,  
3     sendResultByteArrayChannel: SendChannel<ByteArray>,  
4 ) {  
5     val urls = Channel<Url>(capacity = 1)  
6     val data = Channel<DownloadedData>(capacity = 1)  
7     repeat(N_WORKERS) { downloadWorker(urls, data) }  
8     downloader(receiveStringUrlChannel, sendResultByteArrayChannel, urls, data)  
9 }
```

Images downloader - Jetpack Compose demo

```
1 fun main() = application {  
2     val scope = rememberCoroutineScope()  
3     val urlChannel = Channel<String>(capacity = 1)  
4     val imageDataChannel = Channel<ByteArray>(capacity = 1)  
5     LaunchedEffect(Unit) { processUrls(urlChannel, imageDataChannel) }  
6     // application ui
```

Still confused? Check out these places

- 1 Roman Elizarov - Structured concurrency [▶ Watch](#)
- 2 KotlinConf 2017 - Introduction to Coroutines by Roman Elizarov [▶ Watch](#)
- 3 KotlinConf 2017 - Deep Dive into Coroutines on JVM by Roman Elizarov [▶ Watch](#)
- 4 KotlinConf 2018 - Kotlin Coroutines in Practice by Roman Elizarov [▶ Watch](#)
- 5 KotlinConf 2019 - Asynchronous Data Streams with Kotlin Flow by Roman Elizarov [▶ Watch](#)
- 6 Notes on structured concurrency, or: Go statement considered harmful [▶ Read](#)
- 7 Revisiting Coroutines, Ana Moura & Roberto Ierusalimsky [▶ Read](#)



Time to suspend

Thank you for your attention!

